

# COS 417 Precept 5

Traps + Filesystem

# Why Traps?

- The assignment interacts with trapframes in 2 ways:
  - Lazy allocation
  - Copy-on-write
- Traps are the middle ground for syscalls
  - Every syscall is triggered by a trap

Further reading: <https://github.com/palladian1/xv6-annotated/blob/main/traps.md>

# Anatomy of a Trap

- Trap is triggered
  - Interrupt
  - Page Fault
  - Syscall
- Hardware redirects trap to IDT
  - Interrupt Descriptor table, specified in vectors.S
- IDT redirects to trap.asm
- Trapasm.asm sets up trapframe
  - Calls trap in trap.c

```
# vectors.S sends all traps here.
.globl alltraps
alltraps:
    # Build trap frame.
    pushl %ds
    pushl %es
    pushl %fs
    pushl %gs
    pushal

    # Set up data segments.
    movw $(SEG_KDATA<<3), %ax
    movw %ax, %ds
    movw %ax, %es

    # Call trap(tf), where tf=%esp
    pushl %esp
    call trap
    addl $4, %esp
```

# Trap Handler

- `trap.c` handles various cases
  - Syscall? `Trap.c` calls `syscall()`
  - Unknown? `Trap.c` panics/kills process after recording unrecognized input.
  - Page fault? Assignment problem.

# Trap Calls

- Important funcs:
  - lcr3(int) flushes the TLB buffer, updates changes to recent memory
  - rcr2() obtains the relevant memory address in a page fault
- Important trapframe vars:
  - tf->cs&3 returns 0 if trap triggered in kernel mode, 1 otherwise.

```
static inline uint
rcr2(void)
{
    uint val;
    asm volatile("movl %%cr2,%0" : "=r" (val));
    return val;
}

static inline void
lcr3(uint val)
{
    asm volatile("movl %0,%%cr3" : : "r" (val));
}
```

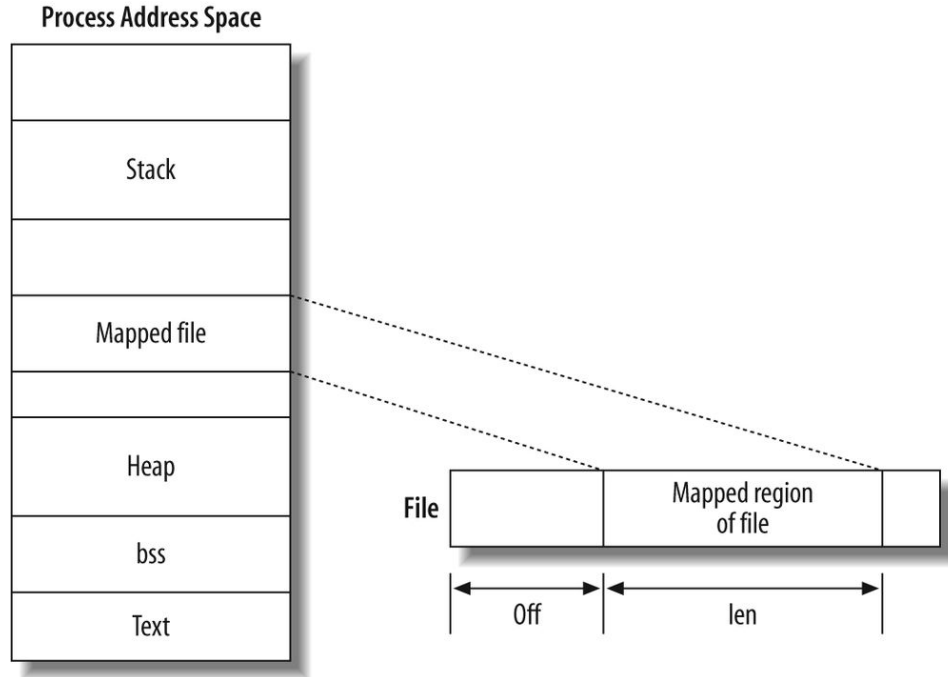
# Filesystems!

- One major feature of wmap is capability of writing to file.
  - Need to learn file systems to support this feature
- You will work directly with file systems later in the course.

# Purpose of file system

- One major feature of wmap is capability of writing to file.
  - Need to learn file systems to support this feature
- You will work directly with file systems later in the course.

# File-based Wmap






# Purpose of file-based wmaps

- Abstraction:
  - Read/Write to file as memory
  - wmap takes care of the rest.
- Performance:
  - Can read from portions large files, without having to seek every time.
  - Multiple processes can access same memory.


# WRT to Assignment

- Implement (limited) version of file-based mapping.

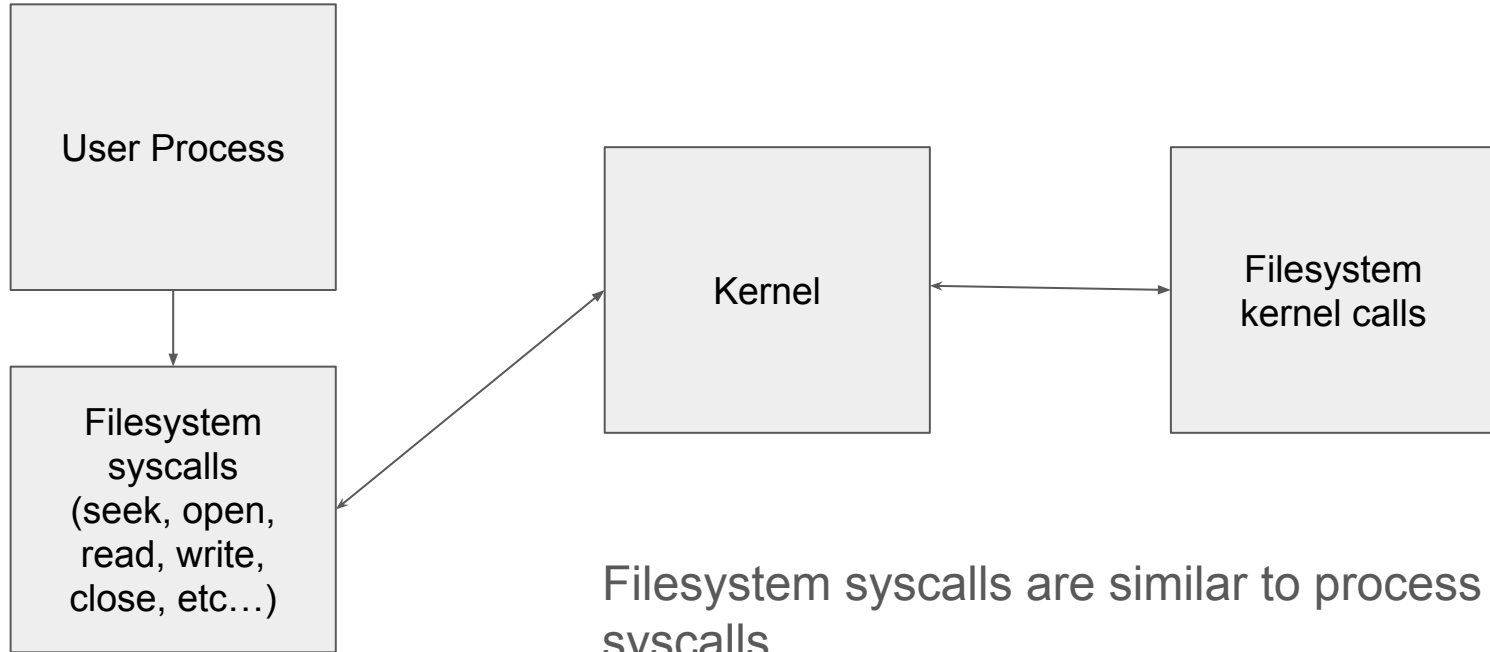


# WRT to Assignment

- Implement (limited) version of file-based mapping.



# Xv6 Filesystem Intro




# Xv6 Filesystem API

Just like process syscalls, filesystem syscalls are wrappers to enter into kernel mode.

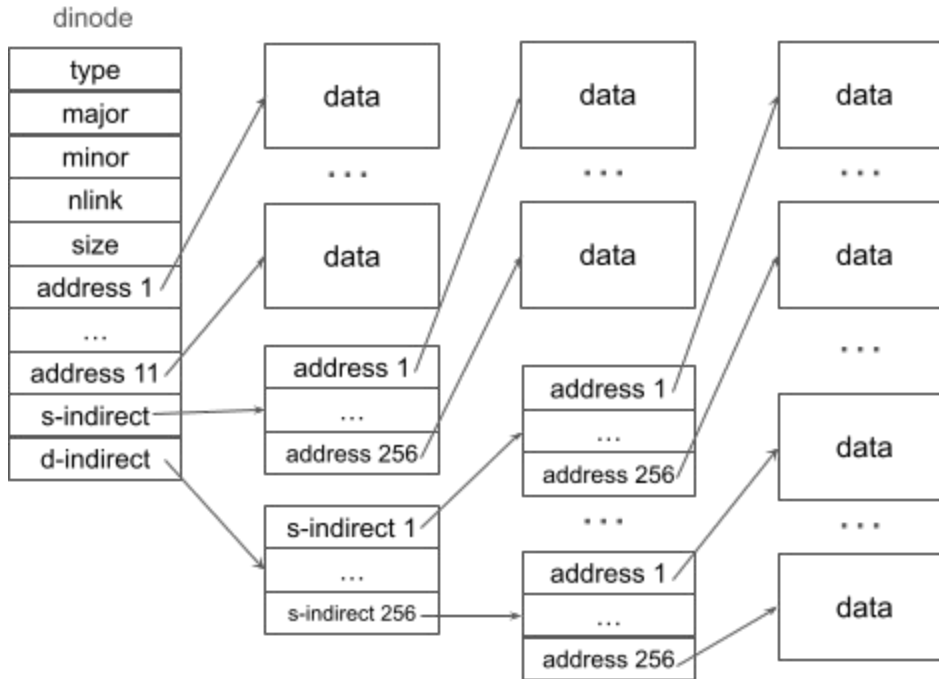
- `syscall()` redirects the syscall into arch.
- For each user function, there exists an equivalent kernel function.
  - Read vs `sysread`,
  - Write vs `syswrite`
  - Etc.

# Xv6 Filesystem API



Calls abstract away from  
filesystem internals

# Xv6 Filesystem Internals



No need to understand this. It's to show a peek behind the hood.

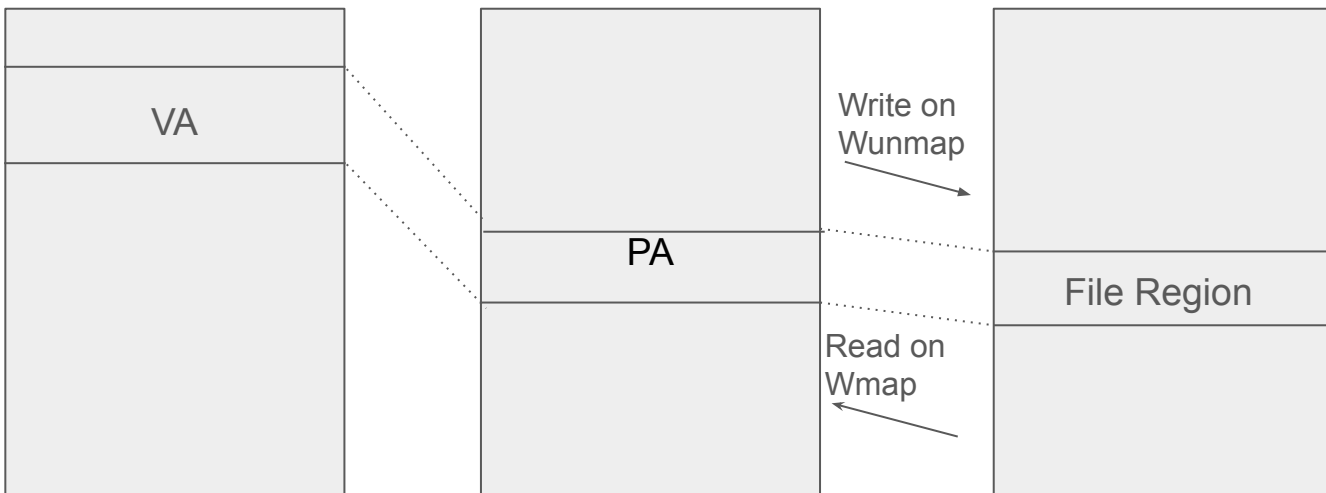
# Xv6 Filesystem Demo

Time to look at the code directly.



# File Wmap in practice

- Implement (limited) version of file-based mapping.



# Summary

- OS handles exceptions/page faults/syscalls at the trap level
- File-based wmap is useful for both performance and convenience reasons.
- For assignment purposes, file-based wmap reads from files at wmap, and writes to files at wunmap.
  - Implications to killing processes.